

Study Report: Learning Spark - Data Engineer

Written by Gagan Pasupuleti

Family: Data Engineer
Level: Intermediate
Chapters: 7
Source: Reference: Learning Spark - Jules Damji et al.

Official sources and free libraries

- <https://www.oreilly.com/library/view/learning-spark-2nd/9781492050045/>

Summary

Study report for *Learning Spark* by Jules Damji et al. (Intermediate level) mapped to the Data Engineer role. PySpark DataFrames, SQL, and cluster processing at scale.

Chapters

Track: Book-Intro

Chapter 1: About Learning Spark [Intermediate]

Chapter focus: About Learning Spark *(Level: Intermediate)**

This study report summarizes *Learning Spark* by Jules Damji et al. for the Data Engineer role. The resource is rated Intermediate level. PySpark DataFrames, SQL, and cluster processing at scale. Use this report alongside the official material to prepare for CodeQuest review.

Code Reference:

```
# Learning Spark
# Author: Jules Damji et al.
# Role: Data Engineer
# Level: Intermediate
```

What it shows: Metadata block records the source book and target role family.

Topic guide

What it actually is

A structured study report based on Learning Spark.

When to use it

When learning Data Engineer skills at Intermediate level.

Where to use it

Data Engineer training paths and certification prep.

Real use example

A learner reads this report before starting the full book or free online guide.

Track: Book-Context

Chapter 2: Why This Book Matters for Your Role [Intermediate]

Chapter focus: Why This Book Matters for Your Role *(Level: Intermediate)**

Data Engineer professionals use ideas from Learning Spark to solve real workplace problems. PySpark DataFrames, SQL, and cluster processing at scale. This chapter explains how the book fits into your learning path and what you should be able to do after studying it.

Code Reference:

Role: Data Engineer

Book focus: PySpark DataFrames, SQL, and cluster processing at scale.

Recommended level: Intermediate

What it shows: Connects the book topic to the job role outcomes.

Topic guide

What it actually is

Role-aligned learning connects theory to job tasks.

When to use it

During career planning and syllabus design.

Where to use it

Data Engineer bootcamps and CodeQuest teacher assignments.

Real use example

A teacher assigns this book report before a intermediate skills checkpoint.

Track: Book-Topics

Chapter 3: Key Topics Covered [Intermediate]

Chapter focus: Key Topics Covered *(Level: Intermediate)**

The main topics in Learning Spark include practical concepts described as: PySpark DataFrames, SQL, and cluster processing at scale. Study each topic with hands-on practice. Take notes on definitions, workflows, and examples that match tools used in Data Engineer jobs today.

Code Reference:

- Topic 1: core concept
- Topic 2: core concept
- Topic 3: core concept
- Topic 4: core concept

What it shows: Topic list guides what to study chapter-by-chapter in the source.

Topic guide

What it actually is

Topic maps turn a book into actionable learning objectives.

When to use it

Before reading and while building chapter summaries.

Where to use it

Self-study, flipped classroom, and revision.

Real use example

A student maps each book chapter to a CodeQuest report section.

Track: Book-Applied

Chapter 4: Applied: PySpark DataFrames (Spark 3.5+) [Intermediate]

Chapter focus: PySpark DataFrames (Spark 3.5+)** *(Level: Intermediate)

PySpark distributes transforms across clusters. Use DataFrame API (not RDDs); lazy evaluation builds logical plans optimized by Catalyst. Cache judiciously - memory is finite.

Code Reference:

```
from pyspark.sql import functions as F

df = spark.read.parquet('s3://bucket/events/')
agg = df.groupBy('country').agg(F.count('*').alias('events'))
agg.write.mode('overwrite').parquet('s3://bucket/gold/events_by_country/')
```

What it shows: groupBy + write parquet is the canonical silver?gold pattern.

Topic guide

What it actually is

PySpark processes terabyte-scale data on commodity clusters.

When to use it

Heavy transforms that choke single-node pandas.

Where to use it

Lakehouse ETL, log aggregation, and ML preprocessing.

Real use example

Daily 500GB clickstream dedupe finishes in 12 minutes on 8 workers.

Chapter 5: Applied: Apache Airflow DAGs [Intermediate]**Chapter focus: Apache Airflow DAGs** *(Level: Intermediate)**

Airflow DAGs define task dependencies with operators (Python, SQL, S3). Use idempotent tasks; retry with exponential backoff; store secrets in Connections/Variables backend (Vault, AWS SM).

Code Reference:

```
from airflow import DAG
from airflow.operators.python import PythonOperator
from datetime import datetime

with DAG('daily_orders', start_date=datetime(2025,1,1), schedule='0 2 * * *') as dag:
    extract = PythonOperator(task_id='extract', python_callable=extract_orders)
    load = PythonOperator(task_id='load', python_callable=load_to_warehouse)
    extract >> load
```

What it shows: >> sets dependency - load waits for successful extract.

Topic guide**What it actually is**

Airflow orchestrates when and in what order pipeline tasks run.

When to use it

Any multi-step pipeline needing schedules, retries, and observability.

Where to use it

dbt + Spark + quality checks coordinated nightly.

Real use example

Failed load triggers Slack alert; retry succeeds after upstream API blip.

Chapter 6: Applied: dbt Models, Tests, and Documentation [Intermediate]**Chapter focus: dbt Models, Tests, and Documentation** *(Level: Intermediate)**

dbt compiles SQL models with ref() dependencies. Add tests: unique, not_null, relationships. Generate docs site for lineage. Use tags and selectors for CI slim runs.

Code Reference:

```
-- models/fct_orders.sql
select
  order_id,
  customer_id,
  amount
from {{ ref('stg_orders') }}
where status = 'completed'
```

What it shows: ref() builds DAG of models - runs in correct order automatically.

Topic guide

What it actually is

dbt is analytics engineering: transform in warehouse with software practices.

When to use it

Gold layer metrics consumed by BI and analysts.

Where to use it

Snowflake/BigQuery + GitHub Actions CI.

Real use example

relationships test catches orphan customer_id before executives see wrong LTV.

Track: Book-Plan

Chapter 7: Study Plan and Practice [Intermediate]

Chapter focus: Study Plan and Practice** *(Level: Intermediate)

Finish Learning Spark with a weekly plan: read one section, write a summary, complete one exercise, and reflect on how it applies to Data Engineer. If a free edition exists, practice every example. Submit your notes and one mini-project demo for teacher review.

Code Reference:

```
Week 1: Read + notes
Week 2: Exercises
Week 3: Mini project
Week 4: Review + quiz
```

What it shows: A 4-week plan turns reading into demonstrable skill.

Topic guide

What it actually is

Structured study plans improve retention and portfolio outcomes.

When to use it

After finishing the key topics chapters.

Where to use it

CodeQuest end-of-unit assessments.

Real use example

A learner completes the study plan and uploads a intermediate project screenshot.